

Systems programming

4 – Processes introduction

MEEC LEEC MEAer LEAer MEIC-A
João Nuno Silva





Bibliography

- Time-Sharing Computer Systems John McCarthy, MIT - 1964
- McCarthy, John, et al. "THOR: a display based time sharing system." 1967
 - Proceedings of the April 18-20, 1967, spring joint computer conference.
- Bauer, Walter F. "Computer design from the programmer's viewpoint."
 - 1958, eastern joint computer conference: Modern computers: objectives, designs, applications
- Beej's Guide to Interprocess Communication, Brian Hall
 - Chapters 2
- The Linux Programming Interface, Michael Kerrisk
 - Chapter 24



First computers

- Batch processing computers
 - First computers
 - Slow, one task at the time
 - Transistors, magnetic core memory
- 1945 – ENIAC
 - 18,000 vacuum tubes, 7,200 crystal diodes, 1,500 relays, 70,000 resistors, 10,000 capacitors,
 - 150 kW, 27 t
 - 100 word memory, 5000 SUMs per second
- 1957 - IBM 709
 - 100 kW, 1 Ton
 - 32K words memory, 42000 SUMs per second



Batch-processing

- One work at a time
 - Load code, Load data
 - Process
 - Save results
- Repeat
 - Other users
 - Wait
 - No interactivity
 - Results only at end



1958 Time sharing computing

- The central idea here is that each large metropolitan area would have one or more of these super computers.
- The computers would handle a number of problems concurrently.
 - Input-output equipment installed on their own premises
 - would buy time on the computer
 - the same way we buy power and water from utility companies.
- Tasks are scheduled by a supervisor
 - Each work has a priority



1964 John McCarthy view

- Private computing
 - It is done by time sharing a large computer.
 - Each user has a console that is connected to the computer by a wired channel such as a telephone line.
- The consoles are two kinds,
 - cheap console
 - electric typewriter used for input and output.
 - expensive console
 - includes a cathode-ray-tube unit for the computer to display pictures and text

Computing consoles

- <https://www.youtube.com/watch?v=S8lGyMKH7zw>
- <https://www.youtube.com/watch?v=jxkygWI-Wfs>



1964 - Time sharing operation

- Operation of such a system as it appears to the user of the typewriter console.
 - When the user wants service
 - he simply starts typing in a message requesting the service.
 - The computer is always ready to pay attention to any key that he strikes
 - just as the telephone system is always ready for you to lift the receiver off the hook.
 - As soon as the key is depressed on the typewriter
 - a signal is sent to the computer.
 - The effect of the signal is
 - make the computer interrupt the current program after the current instruction
 - jump temporarily to a program that determines what the typewriter wants.



1964 - Time sharing operation

- At any moment some of the typewriters
 - will be inactive
 - some will be in the middle of entering messages into the computer,
 - some will be in the process of typing out characters
 - some will be in a status of wanting programs run.
- This is one example of a time-sharing system.
 - There are a large number of users
 - Each user has his own console,
 - Each user gets service from the computer whenever he desires it
 - Each user has the computer maintain his files for him.




Time sharing requirements

- Large primary memory
- Interruption system to handle errors as well as input and output
- Completely nonstop operation
- Erroneous programs must be prevented from damaging other programs
 - it must interrupt if the program attempts any memory references outside its allotted region,
 - programs that get into endless loops must be prevented from wasting computer time.
- Advances memory management (relocation/virtual)
- Secondary storage enough to maintain users' files.

PDP-1

- Main memory
 - magnetic-core memory
 - 18-bit word size (three six-bit chars)
 - 4096 words as standard main memory
 - $<=>$ 9Kbytes (8bit) 12000 chars
 - upgradable to 65,536 words
- 93458 operations per second
 - (10.7ms/instruction)
- 2,700 transistors and 3,000 diodes
- 730Kg





Time-sharing application

- PDP-1
 - 8K Words
 - 22x4Kw drum memory
- BBN time-sharing system
 - increase the effectiveness of the PDP-1 computer
 - for applications involving man-machine interaction
- Allowing the five users
 - each at his own typewriter to interact with the computer
 - just as if he had a computer all to himself.
 - 4kwords per users
 - Response time < 1s

Time-sharing application

- THOR 1966
 - Control computer based teaching laboratories
 - Text editing
 - General purpose programming
 - Better interaction CRT vs teletype
 - Understand trade-offs
- 20 user programs
- 28 consoles





Time sharing Computers

- Cheaper
 - Transistors
 - More compact
 - Multiple users
 - Single CPU
- 1959 - Thor on PDP-1
 - 4K Words
 - 1.3 M USD – 2024 equiv
 - 1964 – DEC PDP-7
 - 4K words
 - 700 kUSD – 2024 equiv
 - 1957 – IBM 709 (vacuum tubes)
 - 32 kwords
 - 1960 IBM 7090
 - 6x faster 3x cheaper than 709
 - 1964 - IBM S/360 Model 67
 - 512 Kbyte



After time-sharing...

- Multitasking/Interactive Operating Systems
 - multiple-programs
 - Single CPU
 - Word and excel running at the same time
- Parallel computers
 - Multiple processors in the same computer
- Multi-threaded Operating Systems
 - Multiple tasks inside the same program
 - Excel, loading a file while updating a table

Processes




Protection

- Operating Systems central role
 - isolation of misbehaving applications
- Fundamental to other OS goals
 - Reliability
 - Security
 - Privacy
 - fairness
- OS kernel
 - lowest level SW running on the system
 - implements protection


Protection

- Applications
 - untrusted
- Process
 - execution of application program with restricted rights
 - Needs permission
 - from OS kernel
 - to access resources (memory from other processes, I/O. ...)



Process Abstraction

- Process:
 - an instance of a program, running with limited rights
- Address space:
 - set of rights of a process
- Memory that the process can access
- Other permissions the process has
 - e.g., which system calls it can make, what files it can access)





UNIX Process Management


- Basic API to manage processes
 - Creation
 - Termination
 - Reception of return code
 - Remember C exit function



UNIX Process Management

- UNIX fork
 - system call to create a copy of the current process, and start it running
 - No arguments!
- UNIX exit
 - system call to terminate a process
- UNIX wait
 - system call to wait for a process to finish
- UNIX signal
 - system call to send a notification to another process

Unix process creation

- Existing process (father)
 - Calls fork system call
 - Waits for start of new project
 - Continues execution
 - New process (son)
 - Is created
 - Continues execution
 - Kernel
 - Receives request
 - Replicates calling process
 - Starts new process
 - Notifies existing process
- 



Unix process creation


- To start a new process call:
 - `#include <unistd.h>`
 - `pid_t fork();`
- A new process is created
- **Man fork**

Unix fork

- The new instruction to be executed is the one following the fork
 - On the father and son !!!!
- All variables are duplicated with the same value
 - After the fork all changes in values are local to each process.
- Fork is a system call that return two different values:
 - In the son return 0
 - In the father return the son PID

Unix fork

```
int new_pid = fork();  
if (new_pid == 0) {  
    printf("child PID #%d\n",  
          getpid());  
    exit(0);  
} else {  
    printf("Parent PID %d\n",  
          child_pid);  
    exit(0);  
}
```



Unix Process creation policies

- Execution mode
 - Father and son execute concurrently
- Memory management
 - Son gets a NON SHARED copy of father memory
- Resource sharing (files, ...)
 - Father and son share all resources
- Child inherits
 - Copy of most information
 - Copy of Memory space
 - Variables
 - Allocated memory
 - Code
 - Opened files
 - pipes, fifos, sockets



Process termination

- When a process executes the exit(int) function:
 - The calling process is terminated "immediately".
 - Any open file belonging to the process are closed;
 - Process's parent is sent a SIGCHLD signal.
- **Man 2 exit / man 3 exit**
 - What happens to the return code?



Reception of return code

- If parent needs the child return code:

- It must wait for its dead:
 - Call **wait /waitpid** function

pid_t waitpid(pid_t pid, int *status, int options);

- Process waits for a specific child termination.
 - 1st argument – ID of child (-1 any process)
 - 2nd argument – child status
 - 3rd argument
 - WNOHANG: return immediately if no child has exited
 - WUNTRACED: also return if a child has stopped

Reception of return code

- **WIFEXITED(status)**
 - returns true if the child terminated normally, exit(3) _exit(2) or by returning from main().
- **WEXITSTATUS(status)**
 - returns the exit status of the child.
- **WIFSIGNALED(status)**
 - returns true if the child process was terminated by a signal.
- **WTERMSIG(status)**
 - returns the number of the signal that caused the child process to terminate.
- **WIFSTOPPED(status)**
 - returns true if the child process was stopped by delivery of a signal
- **WSTOPSIG(status)**
 - returns the number of the signal which caused the child to stop.

Next on PSIS

- Pipes